

First Assignment

Abhijit Chowdhary
High Performance Computing
Prof. Stadler

February 11, 2019

Before we begin, I am running the code Intel Core i5-3320M CPU @ 2.60GHz. This is a processor with 2 cores and 4 threads (on a thinkpad X230).

1 Matrix-Matrix Multiplication

Here is the reported Glop/s and GB/s for various compiler flags.

(a) This is with $-O0$.

Dimension	Time	Gflop/s	GB/s
20	0.005186	0.451190	0.617012
40	0.024997	0.758485	1.024115
60	0.073145	0.878534	1.181222
80	0.177361	0.860618	1.154708
100	0.352622	0.846516	1.134360
120	0.616006	0.838044	1.122067
140	1.004585	0.816516	1.092591
160	1.492329	0.820838	1.097882
180	1.994492	0.874779	1.169621
200	2.748728	0.870948	1.164175
220	3.797066	0.839369	1.121708
240	5.017759	0.824782	1.102006
260	6.099069	0.862863	1.152700
280	8.050332	0.816592	1.090738
300	9.613858	0.841130	1.123378
320	12.707382	0.772389	1.031463
340	13.980333	0.842173	1.124551
360	16.570923	0.843487	1.126214
380	19.782859	0.831019	1.109486
400	23.500127	0.815996	1.089356
420	26.449047	0.839347	1.120464

440	30.684599	0.831888	1.110446
460	35.834093	0.814003	1.086518
480	42.545943	0.778994	1.039742
500	48.110461	0.778677	1.039275
520	55.696070	0.756639	1.009824
540	64.705882	0.729384	0.973414
560	74.624145	0.705372	0.941336
580	84.678332	0.690651	0.921662

(b) This is with $-O3$.

Dimension	Time	Gflop/s	GB/s
20	0.000471	4.971076	6.798053
40	0.004728	4.009766	5.414031
60	0.008584	7.485864	10.065028
80	0.018819	8.111148	10.882882
100	0.039227	7.609594	10.197111
120	0.069211	7.458908	9.986822
140	0.108867	7.534488	10.081991
160	0.189460	6.465538	8.647742
180	0.234924	7.426823	9.930014
200	0.329532	7.264849	9.710742
220	0.429543	7.419845	9.915663
240	0.574186	7.207700	9.630330
260	0.707709	7.436195	9.934030
280	0.902064	7.287552	9.734119
300	1.123595	7.196988	9.612004
320	1.744750	5.625470	7.512364
340	1.646584	7.150475	9.548008
360	1.981027	7.055612	9.420566
380	2.288765	7.182888	9.589802
400	3.252054	5.896580	7.871947
420	3.168978	7.005395	9.351660
440	3.664246	6.966278	9.298938
460	4.171241	6.992897	9.334009
480	5.880642	5.635956	7.522444
500	6.201603	6.040777	8.062432
520	11.679547	3.608174	4.815529
540	24.118377	1.956826	2.611519
560	40.962487	1.285024	1.714896
580	60.356087	0.968968	1.293073

These were computed with the changes:

```
double flops = 3*n*m*(2*k-1)*NREPEATS/ time / 1e9;
double bandwidth = NREPEATS*n*m*k*sizeof(double) / time / 1e9;
```

2 Laplace equation in one space dimension

The error metric I used is the Euclidean Norm. For part (c)

1. Choosing $N = 100$.
 - (a) For Gauss-Seidel, I found that it took 14176 iterations to take the initial error of $10 \rightarrow 10(1e - 6)$ in about .12 seconds. In addition it took about 2994 to go from an initial error of $10 \rightarrow .5$ in about .0027 seconds.
 - (b) For Jacobi, I found that it took 28348 iterations to take the initial error of $10 \rightarrow 10(1e - 6)$ in about .1 seconds. In addition it took about 5985 to go from an initial error of $10 \rightarrow .5$ in about .0044 seconds.
2. Choosing $N = 10,000$, I wasn't able to get the error to come down from $100 \rightarrow 100*(1e-6)$ since it stopped getting finer at around .07. Instead I computed the error going down to .5. My results were:

Jacobi Laplace Solve for $N = 10000$ and 105260486 iterations: 3393.516600

Gauss-Seidel Laplace Solve for $N = 10000$ and 52630244 iterations: 3950.726431s.

Regarding running $N = 10,000$ for 100 iterations using `-O0` and `-O3`: For `-O0` Jacobi took .026 seconds and Jacobi took .027 seconds; for `-O3` Jacobi took .0055 seconds and Gauss-Seidel took .011 seconds. It seems like the optimization is doing some work on the Jacobi method. Included are the relevant bits of the code.

```
double ResidLaplace(long N, double *u, double *f)
{ //Compute ||Au - f|| under the Euclidean Norm
    double hsq = 1.0/((N+1)*(N+1)), a_diag = (2.0/(hsq)), a_bdiag = -1.0/hsq;
    double v = 0.0, tmp = 0.0;

    tmp = u[0]*a_diag + u[1]*a_bdiag - f[0]; v += tmp*tmp;
    tmp = ( u[N-1]*a_diag + u[N-2]*a_bdiag ) - f[N-1]; v += tmp*tmp;
    for (int i = 1; i < N-1; i++)
    {
        tmp = ( a_diag*u[i] + a_bdiag*(u[i-1] + u[i+1]) ) - f[i];
        v += tmp*tmp;
    }
    return sqrt(v);
}

/* Numerically approximates the solution to:
 * -u'' = f in (0,1) with u(0) = 0 and u(1) = 0
 * using the Jacobi method.
 * INPUT: long N: Number of discretization points
 *         double tol: Tolerance to compute solution to.
```

```

*      double *f: discretization of f with N points
*      double *u0: Initial guess vector for u
* OUTPUT: double *u: Array of size N holding solution to equation.
*      long it: Number of iterations needed to hit tol.
*/
long Laplace1DJacobi(long N, double tol, double *f, double *u0, double *u)
{ //Proceed via Jacobi iteration:
    double hsq = 1.0/((N+1)*(N+1)), a_diag = (2.0/(hsq)), a_bdiag = -1.0/hsq;
    double nu[N];
    for (long k = 0; k < N; k++) { u[k] = u0[k]; nu[k] = u0[k]; }

    //Compute  $u_i^{k+1} = 1/a_{ii} (f_i - \sum_{j \neq i} a_{ij} u_j^k$ 
    long it = 0;
    double cur_res = 0.0;
    while (ResidLaplace(N, u, f) > tol && it < 100)
    { //Note for laplace, the summation only has 2 terms max.
        nu[0] = (1/a_diag)*(f[0] - a_bdiag*u[1]);
        nu[N-1] = (1/a_diag)*(f[N-1] - a_bdiag*u[N-2]);
        for (long i = 1; i < N-1; i++)
        {
            nu[i] = (1/a_diag)*( f[i] - a_bdiag*(u[i-1] + u[i+1]) );
        }
        for (long j = 0; j < N; j++) { u[j] = nu[j]; }
        cur_res = ResidLaplace(N,u,f);
        it = it + 1;
    }
    return it;
}

/* Numerically approximates the solution to:
* -u'' = f in (0,1) with u(0) = 0 and u(1) = 0
* using the Gauss Siedel method.
* INPUT: long N: Number of discretization points
*      double tol: Tolerance to compute solution to.
*      double *f: discretization of f with N points
*      double *u0: Initial guess vector for u
* OUTPUT: double *u: Array of size N holding solution to equation.
*      long it: Number of iterations needed to hit tol.
*/
long Laplace1DGaussSeidel(long N, double tol, double *f, double *u0, double *u)
{ //Proceed via Seidel iteration:
    double hsq = 1.0/((N+1)*(N+1)), a_diag = (2.0/(hsq)), a_bdiag = -1.0/hsq;
    double nu[N];
    for (long k = 0; k < N; k++) { u[k] = u0[k]; nu[k] = u0[k]; }

```

```

long it = 0;
double cur_res = 0.0;
while (ResidLaplace(N, u, f) > tol && it < 100)
{ //Note for laplace, the summation only has 2 terms max.
    nu[0] = (1/a_diag)*(f[0] - a_bdiag*u[1]);
    nu[N-1] = (1/a_diag)*(f[N-1] - a_bdiag*nu[N-2]);
    for (long i = 1; i < N-1; i++)
    {
        nu[i] = (1/a_diag)*( f[i] - a_bdiag*(nu[i-1] + u[i+1]) );
    }
    for (long j = 0; j < N; j++) { u[j] = nu[j]; }
    cur_res = ResidLaplace(N,u,f);
    it = it + 1;
}
return it;
}

```